

# COMP9313 2018s2 Project 1 (25 marks)

## Problem statement:

Build the inverted index for a given set of documents (compute the term weights by TF-IDF as shown in slide 5 of Chapter 4, using base 10 logarithm). Ignore the letter case, i.e., consider all words as lower case.

### Input files:

Each line is in format of “DocID DOC”, where DocID is the ID of the document, and DOC is the document content (a list of terms). For example:

0 Construct Inverted Index
1 using MapReduce
2 inverted index code
3 index index file

This sample file “tiny-doc.txt” can be downloaded at:

<https://webcms3.cse.unsw.edu.au/COMP9313/18s2/resources/20465>

### Output:

Set the number of reducers to 3. If a term W is contained in N documents, the corresponding output for this term contains N lines, in format of: “W\tDocID<sub>1</sub>, weight<sub>1</sub>”, “W\tDocID<sub>2</sub>, weight<sub>2</sub>”, ..., and “W\tDocID<sub>n</sub>, weight<sub>n</sub>”. The DocIDs are sorted in ascending order, and the term weights are of double precision (stored in double/DoubleWritable). Given the example graph, the output file is like (assumed in one file):

code\t2, 0.6020599913279624
construct\t0, 0.6020599913279624
file\t3, 0.6020599913279624
index\t0, 0.12493873660829993
index\t2, 0.12493873660829993
index\t3, 0.24987747321659987
inverted\t0, 0.3010299956639812
inverted\t2, 0.3010299956639812
mapreduce\t1, 0.6020599913279624
using\t1, 0.6020599913279624

### Code format:

Name your package as “comp9313.proj1” and name your driver class as “Project1.java”. Finally, package all your java files as a zip file with name “InvertedIndex.zip”.

### Compile:

Your java code will be compiled and packaged as a jar file, and we will use the following commands to check the correctness of your solution:

```
$ $HADOOP_HOME/bin/hadoop jar YOURJAR.jar YOURCLASS input output
```

```
$ $HADOOP_HOME/bin/hdfs dfs -cat output/*
```

Please ensure that the code you submit can be compiled and packaged. Any solution that has compilation errors will receive no more than 5 points for the entire assignment.

## Documentation and code readability

Your source code will be inspected and marked based on readability and ease of understanding. The documentation (comments of the codes) in your source code is also important. Below is an indicative marking scheme:

Result correctness: 18
Efficiency: 5
Code structure, Readability, and Documentation: 2

## Submission:

Deadline: Sunday 9th Sep 09:59:59 PM

Log in any CSE server (e.g., williams or wagner), and use the give command below to submit your solutions:

```
$ give cs9313 project1 InvertedIndex.zip
```

Or you can submit through:

<https://cgi.cse.unsw.edu.au/~give/Student/give.php>

If you submit your assignment more than once, the last submission will replace the previous one. To prove successful submission, please take a screenshot as assignment submission instructions show and keep it by

yourself. If you have any problems in submissions, please email to [xuefeng.chen@student.unsw.edu.au](mailto:xuefeng.chen@student.unsw.edu.au).

## **Late submission penalty**

10% reduction of your marks for the 1st day, 30% reduction/day for the following days.

## **Plagiarism:**

The work you submit must be your own work. Submission of work partially or completely derived from any other person or jointly written with any other person is not permitted. The penalties for such an offence may include negative marks, automatic failure of the course and possibly other academic discipline. Assignment submissions will be examined manually.

Relevant scholarship authorities will be informed if students holding scholarships are involved in an incident of plagiarism or other misconduct.

Do not provide or show your assignment work to any other person - apart from the teaching staff of this subject. If you knowingly provide or show your assignment work to another person for any reason, and work derived from it is submitted you may be penalized, even if the work was submitted without your knowledge or consent.